

STATUS OF PRICING: UNDECIDED. This document deliberately proposes no tiers, price points, or metering model — that is the output we are hiring for. Internal exploratory drafts exist and are available on request, but are excluded here to avoid anchoring. **Goal:** a pricing strategy that is easy to understand, captures the different intents customers arrive with, and stays sound against our cost structure.

1 · THE PRODUCT

Baseout is a B2B SaaS utility for **Airtable platform admins**. It (1) **backs up Airtable automatically** — schema, records, attachments via the REST API on monthly/weekly/daily/webhook-"instant" cadences (all technical capabilities of the system), plus automations & interfaces via user-submitted intake; (2) **restores** — base- or table-level, point-in-time, always into new bases/tables; (3) **gives the data an external home** — the customer's own storage (Drive, Dropbox, Box, OneDrive, S3, BYOS) or Baseout-managed storage, optionally with a **live SQL database copy** (SQLite → shared Postgres → dedicated Postgres → customer's own DB); (4) **layers intelligence on top** — schema visualization & change history, base health scoring, documentation tooling, data metrics/alerts, AI-assisted docs, analytics, and (roadmap) governance, MCP/RAG AI access, connectors. V1 is Airtable-only; architecture is built for multi-platform V2 (Notion, HubSpot, Salesforce), each platform its own subscription component under one billing relationship.

The four jobs customers hire it for — different intents, different WTP curves

JOB	DESCRIPTION	WTP CHARACTER
Insurance	Fear of data loss / accidental deletion; someone accountable for recovery	Cheap peace of mind; price-sensitive
Governance	Best practice & compliance require a copy <i>outside</i> the vendor; Airtable's native snapshots don't satisfy this	Budgeted, requirement-driven
Visibility	No native view of schema complexity, change history, or estate health	Professional/consultant value
Data liberation	Data locked in a proprietary format; a SQL copy makes it queryable, portable, BI-connectable	Buying an unlock, not a backup

The architectural fork any packaging must respect

	STATIC BACKUP	DYNAMIC BACKUP
Mechanism	Flat-file export (CSV) per run to a storage destination	Continuous sync into a managed database
Privacy	With customer-owned storage, record data never touches Baseout servers — a marketed differentiator	Customer explicitly opts in to Baseout hosting their data
Cost shape	Re-transfers the full dataset every run — volume × frequency	Expensive initial sync, cheap deltas, ongoing storage at rest
Enables	Backup + restore only	Everything else: SQL, APIs, real-time sync, all intelligence features

Open question for packaging: separate products, modes within a plan, or a ladder? All intelligence features depend on dynamic mode, so the database is simultaneously a **capability unlock**, a **cost driver**, and a **trust decision**.

Differentiation claims (pressure-test these)

Only Airtable-exclusive backup vendor · most storage destinations in market (only OneDrive support) · only BYOS *and* BYODB vendor (competitors moving away) · "never stores your data" static posture · only vendor unlocking Airtable data as a live SQL layer.

2 · CUSTOMERS & MARKET

Primary buyer: the Airtable platform admin — internal IT / RevOps / BizOps owners and independent consultants managing client accounts. Technically capable (OAuth, storage config, SQL); professionally accountable. Secondary (V2): developers on the SQL layer; enterprise IT/compliance.

SEGMENT	PROFILE & WHAT THEY VALUE	NOTES FOR PRICING RESEARCH
Solo / small ops admin	1 org, a few bases; set-and-forget protection; cheap, simple	Bulk of the legacy base skews here
Consultant / agency	Many client orgs; per-client separation; professional artifacts to show clients	Multi-home one subscription or per-client billing? — packaging question
RevOps / BizOps at scale-up	Airtable business-critical; frequency, retention, SQL for BI, alerts, API	Overlaps the "data liberation" buyer
Enterprise IT / compliance	Governance mandates; BYODB (data never leaves), audit trails, SSO, SLA, CSM	Gated by our SOC 2 timeline (\$6)
Legacy On2Air (~200 paying)	Continuity — "static backups work exactly the same"	Anchored to old prices (below)

Price anchor #1 — what our existing customers pay today (On2Air, being replaced by Baseout)

ON2AIR PLAN	PRICE/MO	BASES	RECORDS	ATTACHMENTS	RESTORES	MAX FREQUENCY
Basic (free)	\$0	1	~1,000	25	1	—
Starter	\$9.99	1	50,000	2,500	1/mo	Monthly
Essentials	\$29.99	15	250,000	25,000	1/mo	Weekly
Professional	\$49.99	50	1,000,000	500,000	5/mo	Daily
Premium	\$79.99	250	5,000,000	1,000,000	10/mo	Hourly

Team observations (inputs, not conclusions): flat fee regardless of actual usage; frequency hard-gated by tier; count caps did the limiting; the team believes On2Air was **historically underpriced** and Baseout is substantially more capable. Migration sensitivity is real — these customers chose cheap and simple; grandfathering / bridge / usage mapping are all open.

Other anchors & landscape

Airtable's own seat pricing (Team ~\$20–24, Business ~\$45–54 per seat/mo) sets expectations for "tooling around Airtable"; a customer's Airtable spend proxies their Baseout budget. **Competitors:** Airtable native snapshots (free "do nothing" default — no external copy, no intelligence, no SQL); Sync Inc (defunct); Whalesync (narrowing); Coefficient (spreadsheet sync); DIY scripts. **Distribution:** Airtable Marketplace listing planned; free pre-registration schema visualization (OAuth in, see your ERD before signup) is the conversion hook.

3 · FEATURE CATALOG (PACKAGING RAW MATERIAL)

Each row is a discrete, gateable unit. Dependencies and cost character are engineering facts; **grouping into packages is open** — internal drafts assign features to tiers; treat all assignments as provisional. Unshipped features can't carry a launch promise. ✓ built · 🟡 partial · ✗ not built

FEATURE	DEPENDENCY	MARGINAL COST CHARACTER	STATUS
CORE BACKUP & RESTORE			
Scheduled backup (monthly/weekly/daily) + manual runs; run history & audit reports	—	Transfer compute per run (static re-sends all)	✓
Instant backup (webhook-driven)	Webhook infra	Low steady-state	✗
Destinations: managed R2, Drive, Dropbox, Box, OneDrive ✓ · S3, Frame.io, custom BYOS ✗	OAuth per provider	R2 = our storage \$; BYOS ≈ zero at rest	🟡
Restore (base/table, point-in-time, to new data)	Retained snapshots	Reverse transfer; Airtable writes slow/rate-limited	✗
Retention windows / smart rolling cleanup	—	Retention = persistent storage cost	✗
Automations & interfaces backup (user-submitted intake)	Intake API/forms	Negligible	✗
DATA INTELLIGENCE (ALL REQUIRE DYNAMIC MODE)			
Schema visualization · changelog · health score · docs	Schema-only DB suffices	Negligible	🟡
AI documentation / schema insights	LLM inference	Cents per generation	✗
Record metrics, data changelog, alerts, PII, analytics	Full dynamic DB	DB storage at rest + scan compute	✗
Governance suite (rules, lineage, retention policy, compliance tooling)	Full dynamic DB	Moderate	✗ V2
ACCESS & INTEGRATION (ALL REQUIRE DYNAMIC MODE)			
Inbound API · SQL REST API · direct SQL connection string	Hosted DB	≈ zero per call; DB instance \$ for direct	✗
MCP / RAG / chatbot / vector DB · Zapier/Make · writeback	Inference + infra	Real per-op cost	✗ V2
ENTERPRISE & TRUST SURFACE			
SSO/SAML, 2FA, SOC 2 (in progress), DPAs, support ladder, Enterprise API	—	Human step costs (support, compliance)	🟡

The database ladder inside dynamic mode (each step a real COGS step — \$5): schema-only SQLite → full SQLite (per-GB) → shared Postgres (amortized) → **dedicated Postgres (fixed \$/mo per instance)** → BYODB (≈ zero infra to us; value is control/compliance).

4 · METERABLE DIMENSIONS (CANDIDATE PRICING METRICS)

Everything the system can count, cap, or gate. No recommendation implied by ordering. "Measured today" = plumbing exists in shipped code.

DIMENSION	VALUE-ALIGNED?	COST-ALIGNED?	MEASURED TODAY	ENFORCEABLE TODAY
Records transferred /mo	Strong — estate size × cadence	Weak-moderate (compute)	✓ per-run counts	✗ no billing pipe
Attachment volume (GB) transferred	Strong	Moderate	✓ per-run counts	✗
Backup frequency (monthly→instant)	Very strong — frequency = RPO (tolerable data loss)	Strong on static; weak on dynamic	✓	✓ shipped
File storage at rest (GB)	Moderate	Strong — persistent, accumulates	🟡 per-run sizes	✗
DB storage (GB) + engine class	Strong (SQL unlock)	Strong — engine class is a step function	✗	✗ provisioning unbuilt
Snapshot retention window	Strong for governance buyers	Strong (storage × time)	✗	✗
Workspaces ("Spaces")	Strong for consultants (per-client)	Weak — unless each carries a dedicated DB	✓	🟡
Bases	Strong — estate-size proxy	Weak	✓	✓ caps shipped
Connections	Moderate	Weak	✓	🟡
Seats / team members	Weak-moderate — admin tool, few users	≈ Zero	✓	✗ not enforced
Restores /mo	High perceived value (crisis moment)	Weak	✗ restore itself not built	
API calls / SQL queries	Strong for integration buyers	≈ Zero	✗ services not built	
AI operations	Moderate	Real per-op inference cost	✗	
Platforms connected (V2)	Strong — new estate protected	Moderate	🟡 billing architecture ready	

Composites the data model could also support (options, not recommendations): a [single abstract usage unit](#) unifying several dimensions (one number, flexible; but customers can't predict bills unaided — internal drafts explored this), or [pure feature-tiering with soft caps](#) (what shipped code does today).

USAGE PHYSICS ANY METERED MODEL MUST HANDLE — **(1) The initial-backup spike:** the first backup of a large base can be 100–200× steady-state usage; without grants/amortization the first invoice ambushes the customer. **(2)** Static mode re-transfers everything each run — usage grows with frequency even if data never changes; dynamic inverts this (deltas only). **(3)** Storage persists while activity resets — retention promises accumulate COGS monthly. **(4)** Airtable rate limits bound transfer speed — big backups cost *time*, not just dollars.

5 · COST STRUCTURE

Fully serverless/usage-based stack — near-zero COGS at rest, scales with activity: Cloudflare Workers + Durable Objects, Trigger.dev v3 (per-second compute for backup tasks), DigitalOcean PostgreSQL master DB, Cloudflare D1/R2/KV/Email, Workers AI, Stripe (~2.9% + 30¢), PostHog, dub.co. The Airtable API is free — its rate limits, not dollars, are the constraint.

Marginal cost anchors (early-2026 list prices — verify)

ACTIVITY	ORDER OF MAGNITUDE
Transfer 1,000 records	Fractions of a cent
Transfer 1 GB attachments	~\$0.01/GB or less (zero egress on R2)
Store 1 GB in R2, 1 month	~\$0.015
Store 1 GB in D1 (SQLite), 1 month	~\$0.75 + per-row ops
Shared Postgres instance	~\$15–60/mo ÷ tenants
Dedicated Postgres per workspace	~\$19–60+/mo each — the big step cost
AI doc generation / restore / API call / email	Cents / low \$ / ≈zero / ≈zero

Illustrative cost-to-serve (for floor-setting)

PROFILE	COGS/MO
Small static/BYOS — 1 base, 10K rec, 0.5 GB, weekly to their Drive	< \$0.50
Medium managed — 5 bases, 250K rec, 10 GB, weekly to R2	~\$1–3
Large dynamic-shared — 20 bases, 1M rec, 50 GB, daily incremental, shared PG	~\$5–15
Heavy dedicated — 50 bases, 5M rec, 200 GB, daily + dedicated PG, long retention	~\$30–80+

Marginal cost is tiny at the low end; dominated by [database instances](#) and [storage-retention promises](#) at the high end. Static/BYOS costs ≈ nothing at rest; enterprise BYODB is the [cheapest to serve](#) — its margin is support/compliance, not COGS.

Fixed & step costs: base platform is low hundreds of \$/mo at launch scale. Steps: dedicated DB instances (above), human support (chat / CSM / SLA where promised), SOC 2 certification, payment processing ~3% of revenue. **Asymmetries to exploit:** zero egress makes restore/download nearly free; dynamic incremental sync is cheaper for us *and* the customer at high frequency; retention is the quiet cost accumulator.

6 · CONSTRAINTS & ASSETS

Billing infrastructure (asset)

- Stripe: one subscription per customer org, **one subscription item per platform** — V2 multi-platform slots in without re-architecting.
- Entitlements resolve from Stripe **product metadata**, never names — new packages are config, not code.
- Metered billing, seat quantities, coupons, per-item trials all supported — no known blocker to tiered, usage-based, hybrid, or seat models.
- Data model designed so **plan limits live in database rows, not code** — limit changes without deploys.
- Trial mechanism live: \$0 subscription at signup, no card; currently 7 days + one backup run. Trial shape (length, caps, free-tier-or-not) is **open for you to revisit**.

Migration constraint

~200 On2Air customers must land somewhere. Believed-real constraints: their current prices are the anchor; their product experience must not regress; legacy accounts can be flagged to see new capabilities as upgrade prompts rather than losing anything. Grandfathering duration, bridge pricing, and mapping are open — internal sketches exist if useful.

Enforcement reality (constraint)

Can gate/meter today: backup frequency, bases-per-workspace caps, trial limits, per-run record/attachment counts (recorded, not billed). **Not built yet:** usage metering/billing pipeline, storage quotas, seat enforcement, restore (entire feature), retention/cleanup, dynamic DB provisioning, webhooks/instant, API/SQL services, AI features.

Implication: a launch model can only enforce what's built — if your recommendation depends on metering or unshipped features, please make the sequencing explicit (launch simple → layer metering, or build first).

Timeline & roadmap constraints

- **SOC 2 in progress** (~6-month runway from ~Apr 2026) — enterprise/governance segment can't be broadly marketed until certified.
- GDPR/DPAs required before hosted-data plans launch.
- **V2 roadmap must extend, not break, the model:** per-platform pricing, possible multi-platform discount, AI features with real inference COGS.

Product-experience commitments (soft — challenge if warranted)

- Restore should never feel punitive at the moment of crisis.
- Bill-shock protection (caps, alerts, predictable defaults) is table stakes if any usage meter exists.
- Pricing shouldn't force privacy-motivated customers into hosted storage (the BYOS "we never store your data" posture is marketed).
- Pre-registration schema visualization stays free (conversion hook).

7 · THE DECISION SPACE — WHAT WE'RE ASKING YOU TO DETERMINE

1. Value metric(s). Which dimension(s) from §4 carry price? Single metric, composite unit, pure feature tiers, or hybrid?

3. Price levels. Against the anchors: legacy On2Air prices, Airtable's own seat spend, competitors, our cost floors (§5), and the "historically underpriced" belief.

5. Expansion & overage philosophy. How customers grow spend — upgrades only, metered overage, add-on packs, seats — and the defaults that prevent bill shock.

7. Enterprise definition. What qualifies as custom-priced; the floor; what's reserved (BYODB, SSO, SLA, CSM).

2. Packaging architecture. How many packages; what anchors each (frequency? DB ladder? intelligence? segment identity?); add-on vs. included; how static and dynamic relate commercially.

4. Free motion. Free tier vs. trial-only; trial shape (time, runs, data caps); what the free schema-visualization hook feeds into.

6. Migration mapping. Where each On2Air cohort lands, what's grandfathered, for how long, and the messaging.

8. Sequencing & multi-platform readiness. What launches given enforcement reality (§6), what it evolves into as metering ships, and how the model extends when Notion/HubSpot arrive.

Success criteria (from the founders): easy to understand · captures customers' different intents · sound against the cost structure · gives the ~200 legacy customers a respectful path · leaves room for the V2 roadmap.

8 · DATA & MATERIALS AVAILABLE ON REQUEST

- On2Air customer roster with plan, tenure, and usage shape (bases, records, attachment volumes, frequency) — for cohort/WTP analysis.
- Shipped telemetry: per-run record/table/attachment counts for every backup run.
- Internal exploratory pricing drafts (tiers, a credit-style meter, overage/add-on sketches) — withheld here to avoid anchoring; available as a strawman whenever useful.
- Product spec set (PRD, features spec, credit-system design, implementation plan) and infrastructure billing exports (Cloudflare, DigitalOcean, Trigger.dev, Stripe) to firm up §5 anchors.